



# Performance of Gensim Word2vec

-----detailed version

Minmei Wang  
mwang107@ucsc.edu

<https://github.com/RaRe-Technologies/gensim>

# Outline of Project

## ➤ Throughput benchmark

- Measure **IPC**, **branch MPKI**, **L1 MPKI**, **LLC MPKI** for gensim
- InputSize:  $X=6,678,526,848,423$  (6T),  
 $\sim X/4, \sim X/16, \sim X/64, \sim X/256$
- Threads: 1, 2, 4, 8

## ➤ Latency benchmark

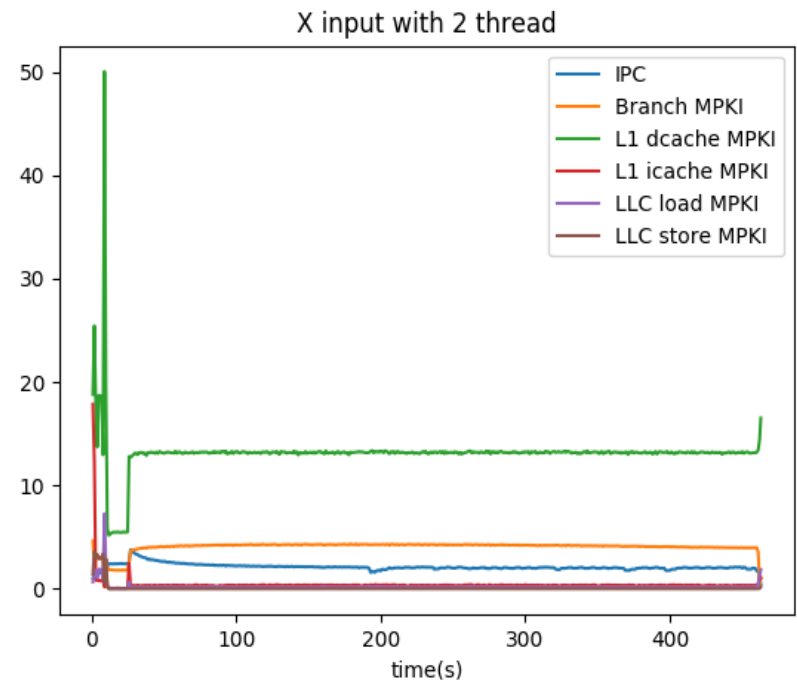
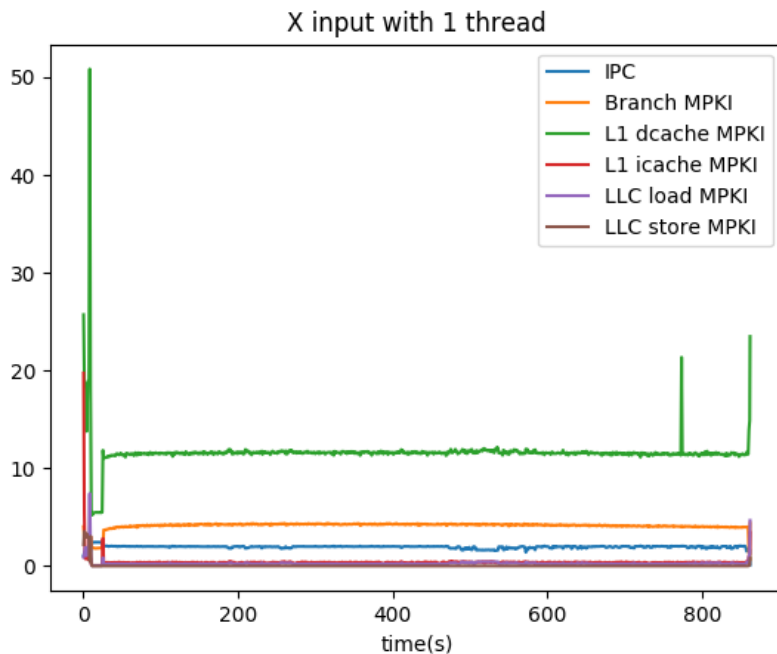
- Measure **latency** for gensim Word2vec
- Total run 200 times for the input

# Part I

## Throughput Benchmark

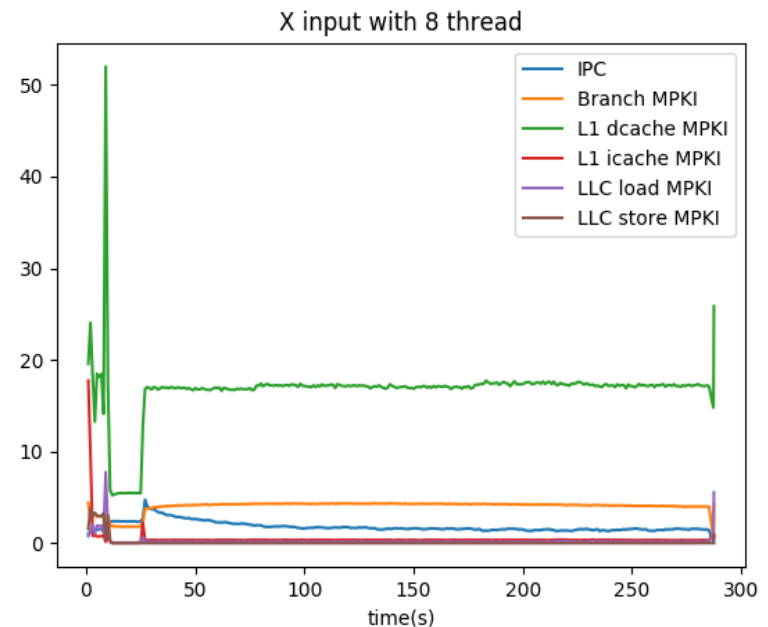
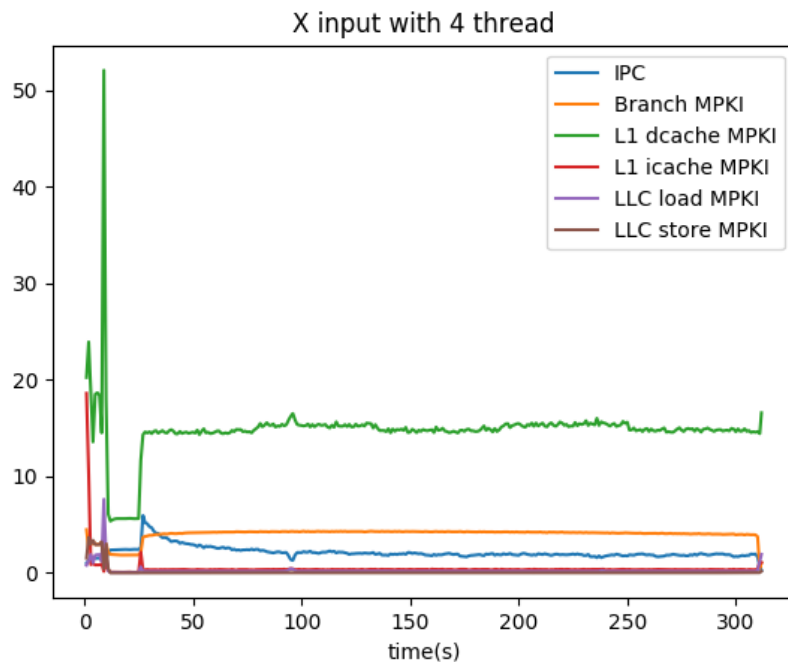
# Throughput benchmark

➤ Input Size: X (6,678,526,848,423) (1,2,4,8 threads)



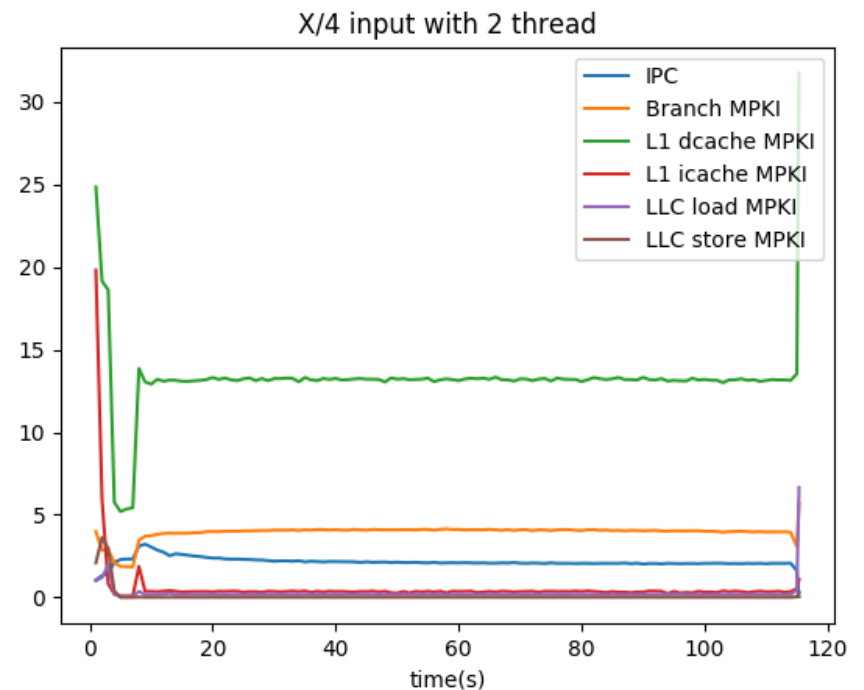
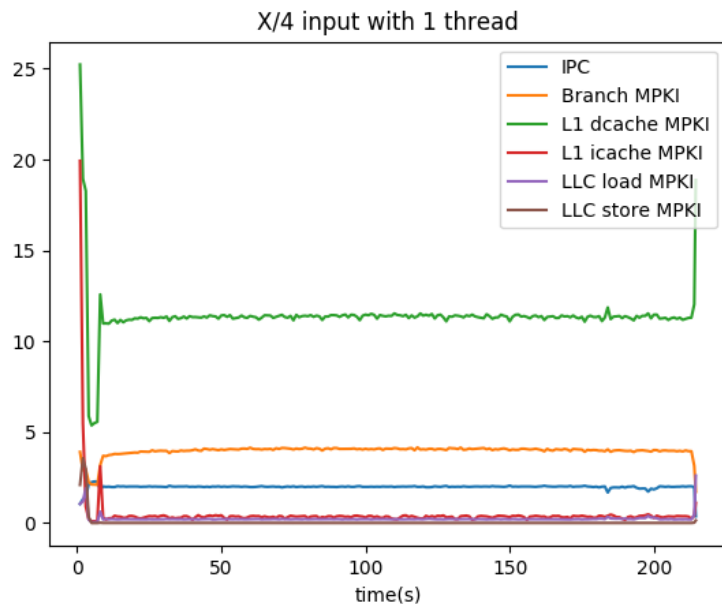
# Throughput benchmark

- Input Size: X (1,2,4,8 threads)
- It has good speedup when threads enlarges to 4.



# Throughput benchmark

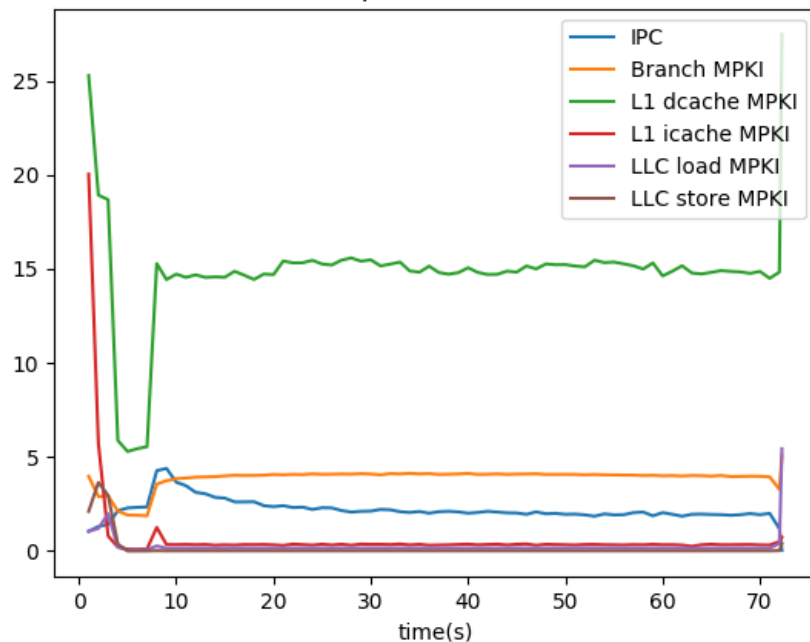
➤ Input Size: X/4 (1,709,851,631,548)(1,2,4,8 threads)



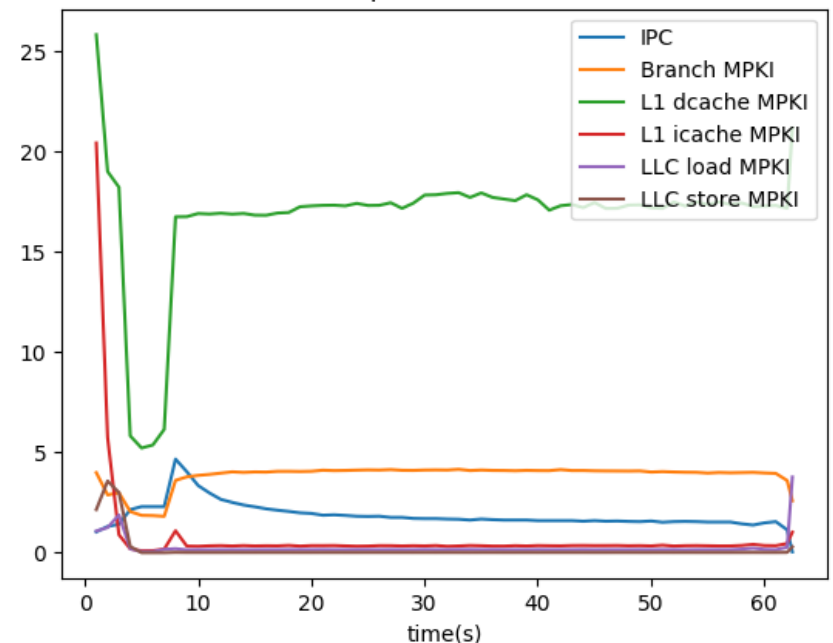
# Throughput benchmark

➤ Input Size: X/4 (1,2,4,8 threads)

X/4 input with 4 thread



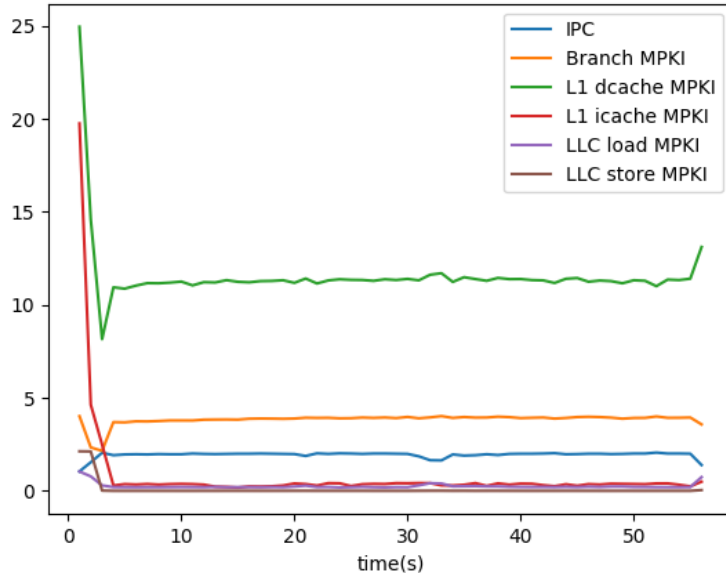
X/4 input with 8 thread



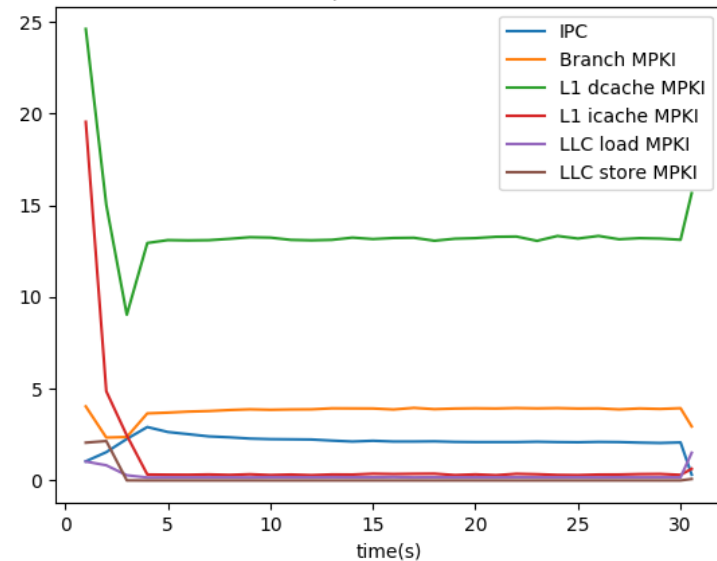
# Throughput benchmark

➤ Input Size: X/16 (437,717,377,815)(1,2,4,8 threads)

X/16 input with 1 thread



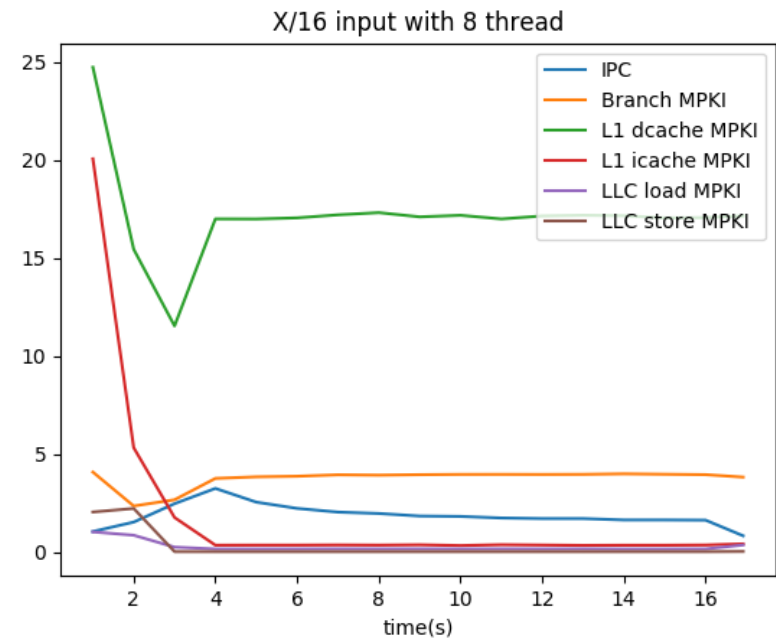
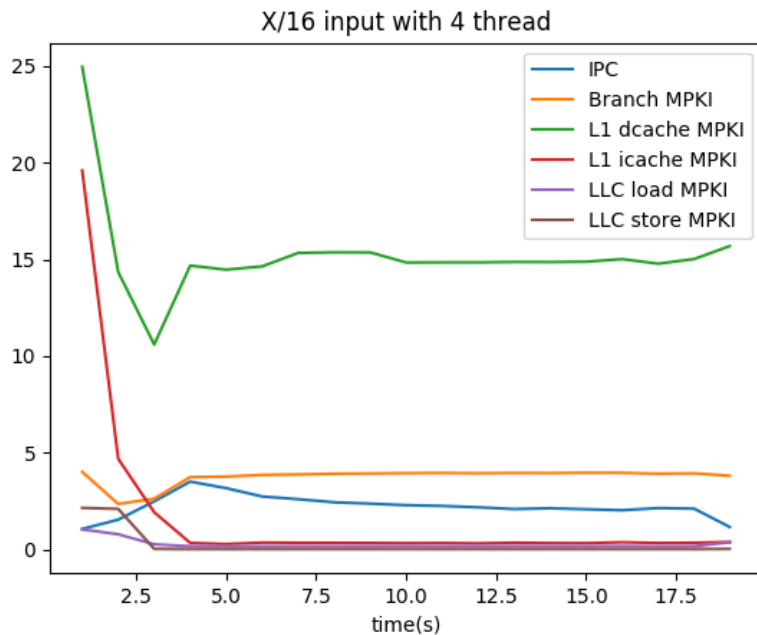
X/16 input with 2 thread





# Throughput benchmark

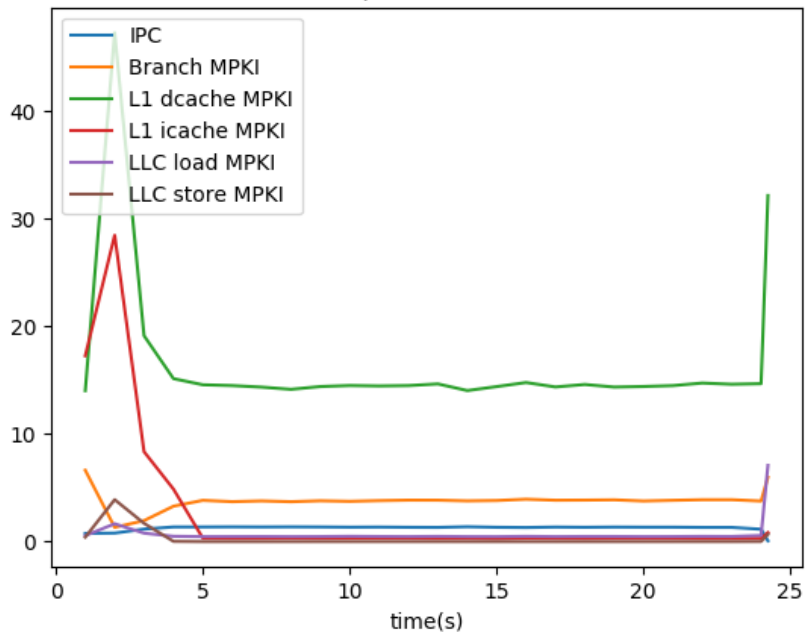
➤ Input Size: X/16 (1,2,4,8 threads)



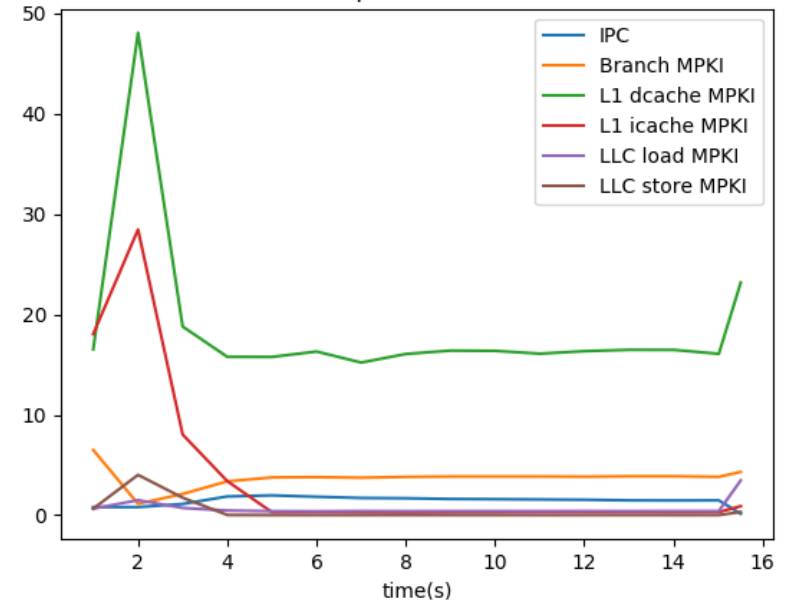
# Throughput benchmark

➤ Input Size: X/64 (113,684,901,010)(1,2,4,8 threads)

X/64 input with 1 thread

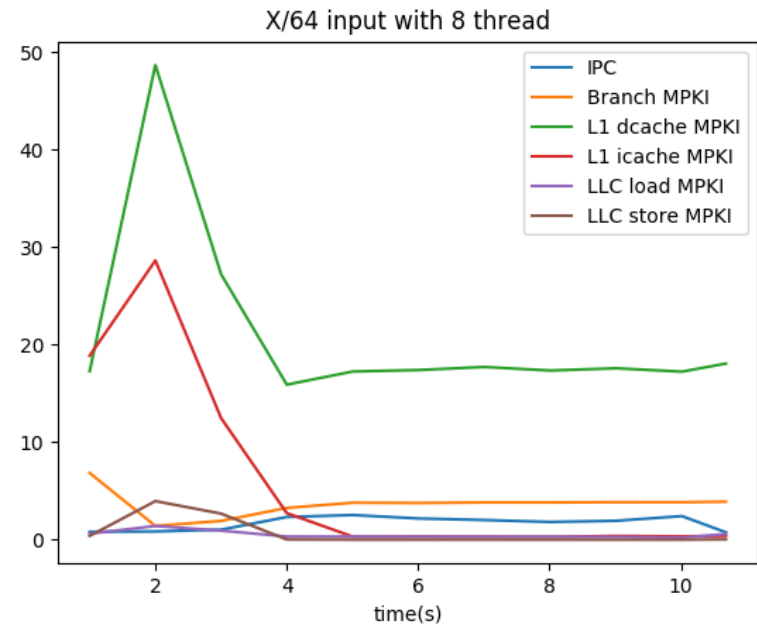
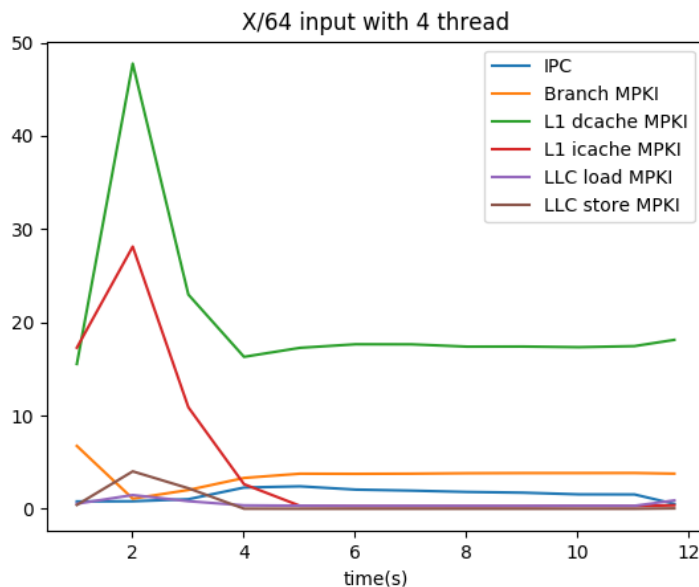


X/64 input with 2 thread



# Throughput benchmark

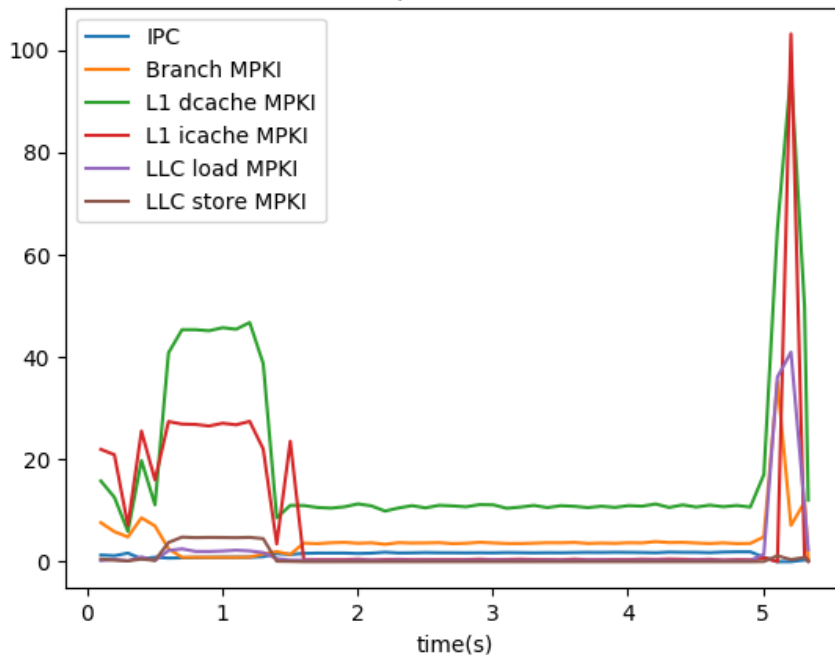
➤ Input Size: X/64 (1,2,4,8 threads)



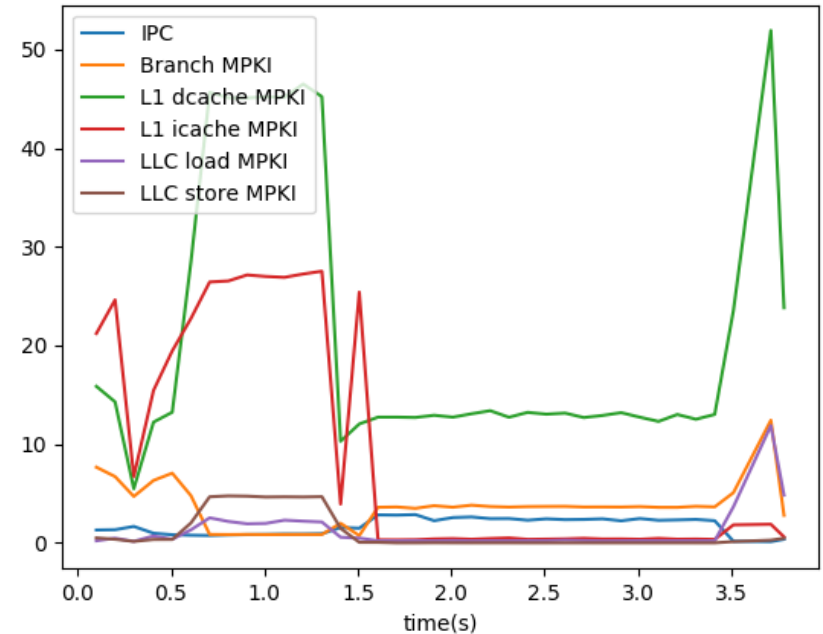
# Throughput benchmark

➤ Input Size: X/256(30,460,388,166) (1,2,4,8 threads)

X/256 input with 1 thread

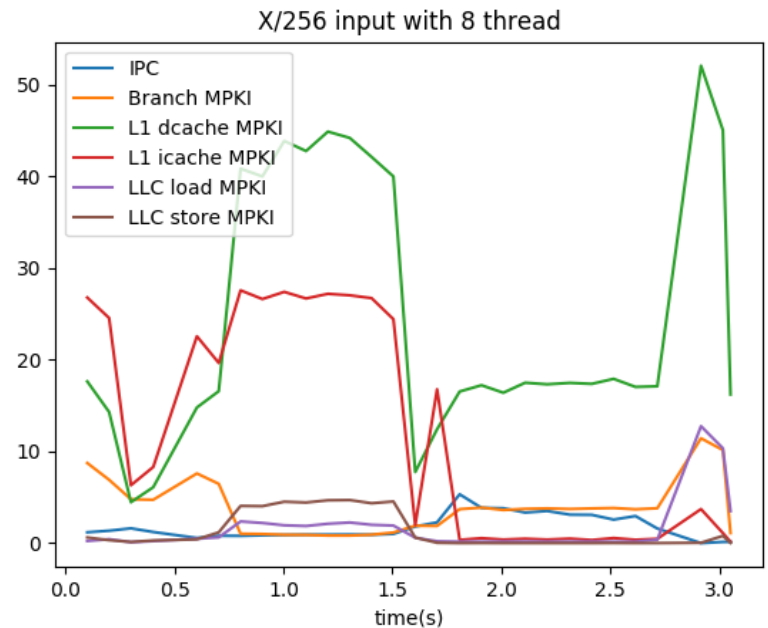
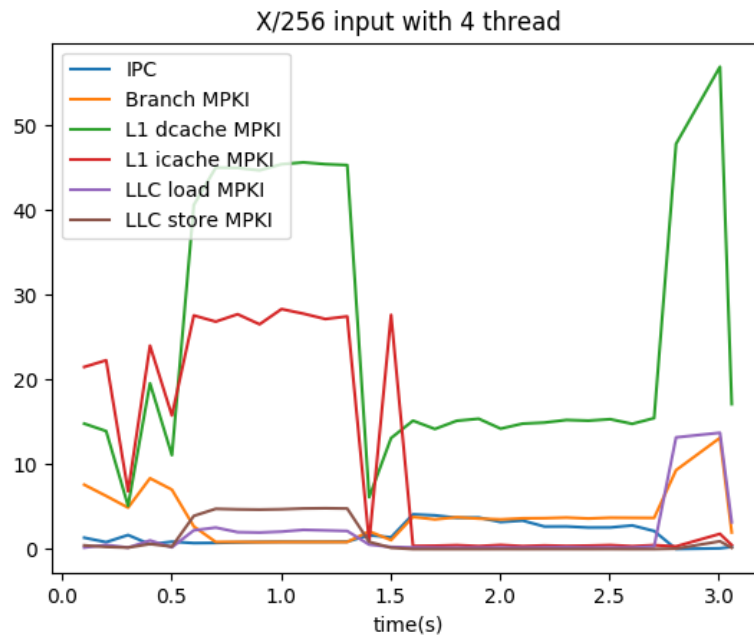


X/256 input with 2 thread



# Throughput benchmark

➤ Input Size: X/256 (1,2,4,8 threads)



# Conclusions

- The runtime is proportional to the size of instructions
- Multithreads
  - **Multithreads** is useful for speed-up, especially when threads is from 1 to 2, 4.
  - It has limitations, It has little speedup when the number of threads is from 4 to 8. It is the reason of data dependencies in the code.
- L1 dcache MPKI is relatively higher in the code
- LLC load and store MPKI performance is good in the code

# Part II

## Latency Benchmark

# Latency Benchmark

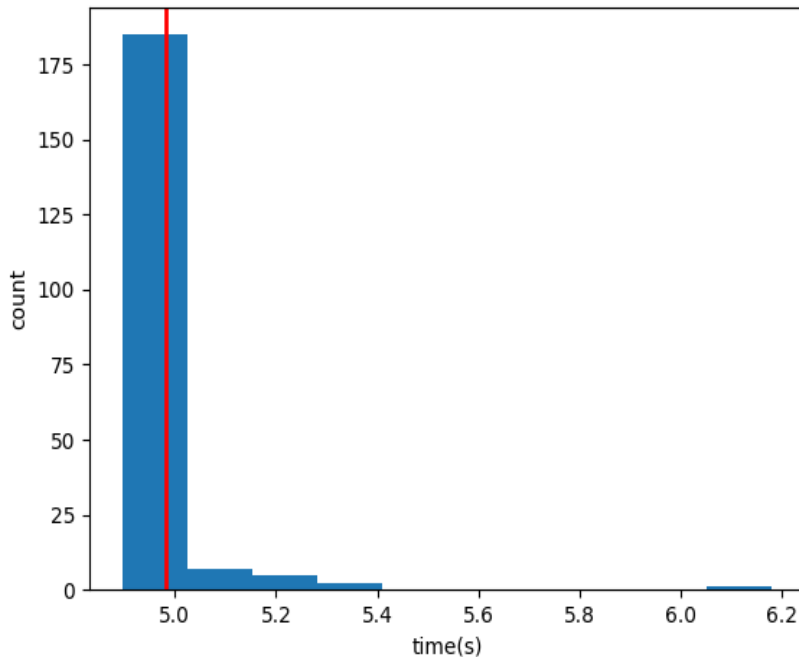
- Latency is the time between giving the input and generating the output
  - Input: A list of files
  - Output: A trained word2vec model



# Latency Benchmark

- I run **200** times for benchmark with almost 30,460,388,166 instructions. The read line is the average.
- The latency distribution with 1 thread is more concentrated may be because the code has strong data dependency during the running process

word2vec latency distribution (1 thread)



word2vec latency distribution (8 thread)

